

# Hardware für Eingebettete Systeme

## Lab 4: AVX intrinsics and masking

Prof. Dr. Mario Porrmann  
M.Sc. Marc Rothmann

# 1 Masking intrinsics

Using the `blend` intrinsics, we can merge two vectors based on a mask that we supply. E.g., the function:

```
__m256 _mm256_blend_ps (__m256 a, __m256 b, const int imm8)
```

Implements the following operation:

```
FOR j := 0 to 7
  i := j*32
  IF imm8[j]
    dst[i+31:i] := b[i+31:i]
  ELSE
    dst[i+31:i] := a[i+31:i]
  FI
ENDFOR
dst[MAX:256] := 0
```

This can be used to control the data flow in AVX programs.  
Implement the following functions using AVX intrinsics:

a) For two vectors  $a$  and  $b$ :

$$c_i = \begin{cases} a_i - b_i & \text{if } a_i \geq b_i \\ a_i & \text{otherwise} \end{cases}$$

b) For a vector  $a$ :

$$c_i = \begin{cases} \text{sqrt}(a_i) & \text{if } a_i \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

*Hint:* It's probably easier to use `_mm256_blendv_ps` instead of `_mm256_blend_ps`.

## 2 Vectorizing K-Means

The K-Means algorithm is a clustering algorithm that aims to partition  $n$  observations into  $k$  clusters. Each cluster is represented by a cluster center (the mean). The observations belong to the cluster with the closest cluster center. K-Means iteratively moves the centers to minimize the within cluster variance.

This is the pseudo-code of the K-Means algorithm:

**Initialize**  $k$  means  $m_0, \dots, m_{k-1}$

**for**  $t = 1, \dots, \text{EPISODES}$ :

1. Assign each data point  $x_j$  to a cluster  $S_i$ :

$$S_i^t = \{x_j : \|x_j - m_i^t\|^2 \leq \|x_j - m_{i^*}^t\|^2 \text{ for all } i^* = 1, \dots, k\}$$


2. Update the mean  $m_i$  of each cluster  $S_i$ :

$$m_i^{t+1} = \frac{1}{|S_i^t|} \sum_{x_j \in S_i^t} x_j$$

**return**  $m_0, \dots, m_{k-1}$

**Algorithm 1:** K-Means

The file `clusters.csv` contains 750 2D-points which can be partitioned into 4 clusters. The following figure visualizes the data.



The cluster centers are:

$$m_0 = (10, 5)$$

$$m_1 = (30, 5)$$

$$m_2 = (20, 15)$$

$$m_3 = (35, 25)$$

- a) The file `kmeans.c` contains un-vectorized k-means code. Implement the vectorized version `kmeans_simd`. You need to use the following compiler flags:

```
gcc -mavx -O3 kmeans.c -lm
```

Verify that your algorithm approximates the correct cluster centers.

- b) Compare the speed of your implementation to the speed of the un-vectorized version. You can use the `eval.py` script to generate a plot showing your computed cluster centers in the data.

**Abgabe zum 4.6.2023!**